

Implement your Own Shell

Theory

- Processes are created with the fork system call
- Fork returns to child process with return value 0.
- The child process created by fork is a copy of the original parent process, except that it has its own process ID.
- When parent executes fork it returns to parent with process id of child .
- This helps parent and child to execute different code after creation of new process.
- After forking a child, parent and child execute simultaneously, this behavior can be modified by using wait() call and making parent wait till execution of child process gets over.
- To execute a command we use system call 'exec'. Varieties of exec are available.
- Pipe() **has to be used** to communicate between 2 processes.

Input

- Get user input using function's such as getline().
- Parse the input line, clean extra white space and store the parameters in an array.

Output Redirection

The UNIX shell provides this nice feature with the ">" character.

If a user types "ls > output" , nothing should be printed on the screen. Instead, the output of the ls program should be re-directed to the output file.

If the "output" file already exists before you run your program, you should simply append at the end of file .

Background Processes

Your shell should able to provide support for background processes. To put a process in the background, the user types & after the command. For example:

```
mysh> ls &
```

```
mysh>
```

When a user backgrounds a process, control returns immediately to the shell, allowing the user to type more. Redirect the standard output for a command that is being run in the background to a file. This prevents the output from the command appearing on your screen and interrupting your current work. Command “**lsb**” should list the currently running background processes.

Inter Process Communication

You could support `|`, a pipe, between two processes. For example, `ps | grep "root"` would send the *STDOUT* of `ps` to the *STDIN* of `grep` using a pipe. You may want to start by supporting pipes only between two processes before considering longer chains. Longer chains will probably require something like handle process n and then recursively handle the other $n-1$. Eg : `ls -l | grep "Aug" | sort +4n`

Functions provided by gcc

- `fork()`
- `exec()`
- `wait()`
- `pipe()`
- `dup2()`
- `chdir()`
- `getcwd()`
- `getlogin()`
- `waitpid`
- `exit`

Program Errors(Boundary Test Cases)

For the following situation, you should print the error message and continue processing:

- A command does not exist or cannot be executed.
- A very long command line (over 512 characters, excluding the carriage return)

Your shell should also be able to handle the following scenarios below, which are not errors. A reasonable way to check if something is not an error is to run the command line in the real Unix shell.

- An empty command line.
- An empty command between two or more `'` characters.
- Multiple white spaces on a command line.
- White space before or after the `'` character or extra white space in general.

```
> ls;ls;ls
> ls ; ls ; ls
> ls>a; ls > b; ls> c; ls >d
```

Simple Shell (Pseudo Code):

```
int main (int argc, char **argv)
{
    while (1){
        int childPid;
        char * cmdLine;

        printPrompt();

        cmdLine= readCommandLine(); //or GNU readline("");

        cmd = parseCommand(cmdLine);

        if ( isBuiltInCommand(cmd)){
            executeBuiltInCommand(cmd);
        } else {
            childPid = fork();
            if (childPid == 0){
                executeCommand(cmd); //calls execvp

            } else {
                if (isBackgroundJob(cmd)){
                    record in list of background jobs
                } else {
                    waitpid (childPid);
                }
            }
        }
    }
}
```

Note : You have to implement your shell using functions listed here. Use of yacc or lex is strictly prohibited.